

Week3_Assignment

February 5, 2024

```
[1]: import pandas as pd
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.svm import SVC
from sklearn.preprocessing import MinMaxScaler
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, classification_report, \
    confusion_matrix, ConfusionMatrixDisplay
import matplotlib.pyplot as plt
```

0.1 Load dataset

```
[55]: df = pd.read_csv("preped_churn_data.csv")

df.sample(10)
```

```
[55]:
```

	tenure	PhoneService	Contract	PaymentMethod	\
2107	33	Yes	Month-to-month	Electronic check	
2138	54	Yes	Two year	Bank transfer (automatic)	
2006	5	Yes	Month-to-month	Electronic check	
5117	66	No	Two year	Bank transfer (automatic)	
1568	11	Yes	Month-to-month	Electronic check	
838	36	Yes	Month-to-month	Credit card (automatic)	
6100	29	Yes	Month-to-month	Bank transfer (automatic)	
4170	72	Yes	Two year	Bank transfer (automatic)	
4209	44	Yes	Month-to-month	Bank transfer (automatic)	
6556	6	Yes	One year	Credit card (automatic)	

	MonthlyCharges	TotalCharges	Churn	\
2107	93.35	3092.00	No	
2138	69.95	3871.85	No	
2006	47.15	223.15	Yes	
5117	61.35	4193.40	No	
1568	111.40	1183.05	No	
838	79.20	2854.95	No	
6100	99.05	2952.85	Yes	
4170	115.15	8349.45	No	
4209	54.90	2549.10	No	

6556	19.00	105.50	No
------	-------	--------	----

	TotalCharges_to_MonthlyCharges_ratio	customerID
2107	33.122657	2108
2138	55.351680	2139
2006	4.732768	2007
5117	68.352078	5118
1568	10.619838	1569
838	36.047348	839
6100	29.811711	6101
4170	72.509336	4171
4209	46.431694	4210
6556	5.552632	6557

0.2 Convert categorical columns to numeric

```
[56]: df['Churn'] = df['Churn'].map({'Yes': 1, 'No': 0})

df['PhoneService'] = df['PhoneService'].map({'Yes': 1, 'No': 0})
df
```

```
[56]:
```

	tenure	PhoneService	Contract	PaymentMethod \
0	1	0	Month-to-month	Electronic check
1	34	1	One year	Mailed check
2	2	1	Month-to-month	Mailed check
3	45	0	One year	Bank transfer (automatic)
4	2	1	Month-to-month	Electronic check
...
7038	24	1	One year	Mailed check
7039	72	1	One year	Credit card (automatic)
7040	11	0	Month-to-month	Electronic check
7041	4	1	Month-to-month	Mailed check
7042	66	1	Two year	Bank transfer (automatic)

	MonthlyCharges	TotalCharges	Churn \
0	29.85	29.85	0
1	56.95	1889.50	0
2	53.85	108.15	1
3	42.30	1840.75	0
4	70.70	151.65	1
...
7038	84.80	1990.50	0
7039	103.20	7362.90	0
7040	29.60	346.45	0
7041	74.40	306.60	1
7042	105.65	6844.50	0

	TotalCharges_to_MonthlyCharges_ratio	customerID
0	1.000000	1
1	33.178227	2
2	2.008357	3
3	43.516548	4
4	2.144979	5
...
7038	23.472877	7039
7039	71.345930	7040
7040	11.704392	7041
7041	4.120968	7042
7042	64.784666	7043

[7043 rows x 9 columns]

0.3 Generating dummies

```
[57]: PM_dummies = pd.get_dummies(df['PaymentMethod'], prefix='PaymentMethod')
      C_dummies = pd.get_dummies(df['Contract'], prefix='Contract')

      df = pd.concat([df, PM_dummies, C_dummies], axis=1)

      df = df.loc[:, ~df.columns.duplicated()]

      df.head(5)
```

```
[57]:   tenure  PhoneService      Contract      PaymentMethod \
0        1             0  Month-to-month      Electronic check
1       34             1      One year      Mailed check
2        2             1  Month-to-month      Mailed check
3       45             0      One year  Bank transfer (automatic)
4        2             1  Month-to-month      Electronic check

      MonthlyCharges  TotalCharges  Churn  TotalCharges_to_MonthlyCharges_ratio \
0          29.85         29.85      0          1.000000
1          56.95        1889.50      0          33.178227
2          53.85         108.15      1          2.008357
3          42.30        1840.75      0          43.516548
4          70.70         151.65      1          2.144979

      customerID  PaymentMethod_Bank transfer (automatic) \
0                1                                   False
1                2                                   False
2                3                                   False
3                4                                    True
4                5                                   False
```

	PaymentMethod_Credit card (automatic)	PaymentMethod_Electronic check \
0	False	True
1	False	False
2	False	False
3	False	False
4	False	True

	PaymentMethod_Mailed check	Contract_Month-to-month	Contract_One year \
0	False	True	False
1	True	False	True
2	True	True	False
3	False	False	True
4	False	True	False

	Contract_Two year
0	False
1	False
2	False
3	False
4	False

```
[65]: dummies = ['Contract_Month-to-month', 'Contract_One year', 'Contract_Two year',
↳ 'PaymentMethod_Bank transfer (automatic)', 'PaymentMethod_Electronic check',
↳ 'PaymentMethod_Mailed check', 'PaymentMethod_Credit card (automatic)']

for column in dummies:
    df[column] = pd.factorize(df[column])[0]

df.sample(5)
```

```
[65]:
```

	tenure	PhoneService	Contract	PaymentMethod \
6261	62	1	Two year	Bank transfer (automatic)
1328	23	1	Month-to-month	Electronic check
6694	8	1	Two year	Credit card (automatic)
4197	22	1	Month-to-month	Electronic check
2169	62	1	Two year	Mailed check

	MonthlyCharges	TotalCharges	Churn \
6261	20.45	1297.35	0
1328	54.25	1221.55	0
6694	76.15	645.80	0
4197	84.75	1816.75	0
2169	81.00	4985.90	0

	TotalCharges_to_MonthlyCharges_ratio	customerID \
6261	63.440098	6262
1328	22.517051	1329

6694	8.480630	6695
4197	21.436578	4198
2169	61.554321	2170

	PaymentMethod_Bank transfer (automatic) \
6261	1
1328	0
6694	0
4197	0
2169	0

	PaymentMethod_Credit card (automatic)	PaymentMethod_Electronic check \
6261	0	1
1328	0	0
6694	1	1
4197	0	0
2169	0	1

	PaymentMethod_Mailed check	Contract_Month-to-month	Contract_One year \
6261	0	1	0
1328	0	0	0
6694	0	1	0
4197	0	0	0
2169	1	1	0

	Contract_Two year
6261	1
1328	0
6694	1
4197	0
2169	1

0.4 Dropping unneeded columns

```
[67]: df = df.drop(['PaymentMethod', 'Contract', 'customerID'], axis=1)
df
```

```
[67]:
```

	tenure	PhoneService	MonthlyCharges	TotalCharges	Churn \
0	1	0	29.85	29.85	0
1	34	1	56.95	1889.50	0
2	2	1	53.85	108.15	1
3	45	0	42.30	1840.75	0
4	2	1	70.70	151.65	1
...
7038	24	1	84.80	1990.50	0
7039	72	1	103.20	7362.90	0
7040	11	0	29.60	346.45	0

7041	4	1	74.40	306.60	1
7042	66	1	105.65	6844.50	0

	TotalCharges_to_MonthlyCharges_ratio \
0	1.000000
1	33.178227
2	2.008357
3	43.516548
4	2.144979
...	...
7038	23.472877
7039	71.345930
7040	11.704392
7041	4.120968
7042	64.784666

	PaymentMethod_Bank transfer (automatic) \
0	0
1	0
2	0
3	1
4	0
...	...
7038	0
7039	0
7040	0
7041	0
7042	1

	PaymentMethod_Credit card (automatic)	PaymentMethod_Electronic check \
0	0	0
1	0	1
2	0	1
3	0	1
4	0	0
...
7038	0	1
7039	1	1
7040	0	0
7041	0	1
7042	0	1

	PaymentMethod_Mailed check	Contract_Month-to-month	Contract_One year \
0	0	0	0
1	1	1	1
2	1	0	0
3	0	1	1

4	0	0	0
...
7038	1	1	1
7039	0	1	1
7040	0	0	0
7041	1	0	0
7042	0	1	0

Contract_Two year	
0	0
1	0
2	0
3	0
4	0
...	...
7038	0
7039	0
7040	0
7041	0
7042	1

[7043 rows x 13 columns]

```
[68]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 13 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   tenure                                7043 non-null   int64
1   PhoneService                         7043 non-null   int64
2   MonthlyCharges                       7043 non-null   float64
3   TotalCharges                         7043 non-null   float64
4   Churn                                7043 non-null   int64
5   TotalCharges_to_MonthlyCharges_ratio 7043 non-null   float64
6   PaymentMethod_Bank transfer (automatic) 7043 non-null   int64
7   PaymentMethod_Credit card (automatic) 7043 non-null   int64
8   PaymentMethod_Electronic check         7043 non-null   int64
9   PaymentMethod_Mailed check             7043 non-null   int64
10  Contract_Month-to-month                7043 non-null   int64
11  Contract_One year                      7043 non-null   int64
12  Contract_Two year                      7043 non-null   int64
dtypes: float64(3), int64(10)
memory usage: 715.4 KB
```

```
[60]: df.isna().sum()
```

```
[60]: tenure          0
      PhoneService    0
      Contract        0
      PaymentMethod    0
      MonthlyCharges    0
      TotalCharges     0
      Churn            0
      TotalCharges_to_MonthlyCharges_ratio  0
      customerID       0
      PaymentMethod_Bank transfer (automatic) 0
      PaymentMethod_Credit card (automatic)  0
      PaymentMethod_Electronic check         0
      PaymentMethod_Mailed check             0
      Contract_Month-to-month                0
      Contract_One year                      0
      Contract_Two year                      0
      dtype: int64
```

0.5 Modelling

0.6 Split data into features (X) and target (y)

```
[69]: X = df.drop('Churn', axis=1)
      y = df['Churn']
```

0.7 Split into training and testing sets

```
[70]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
      ↪random_state=42)

      scaler = MinMaxScaler()
      X_train = scaler.fit_transform(X_train)
      X_test = scaler.transform(X_test)
```

```
[71]: num_rows, num_columns = len(X_train), X_train.shape[1]
      print(f"Number of rows: {num_rows}")
      print(f"Number of columns: {num_columns}")
```

```
Number of rows: 5634
Number of columns: 12
```

```
[72]: y_train.shape
```

```
[72]: (5634,)
```

```
[73]: X_test.shape
```

```
[73]: (1409, 12)
```



```
[74]: y_test.shape
```

```
[74]: (1409,)
```

0.8 Train Model

```
[75]: # Create and train the Linear Regression model
lr_model = LogisticRegression(random_state=42)
lr_model.fit(X_train, y_train)
```

```
[75]: LogisticRegression(random_state=42)
```

```
[76]: print(lr_model.score(X_train, y_train))
print(lr_model.score(X_test, y_test))
```

```
0.7903798367057153
```

```
0.8048261178140526
```

The accuracy score on the training set is 0.79 while the accuracy score on the test set is approximately 0.70, indicating that the model performs less accurately on test data compared to the training data.

0.9 Make predictions on the test data

```
[77]: y_pred = lr_model.predict(X_test)
```

0.10 Evaluation

```
[84]: accuracy = accuracy_score(y_test, y_pred)
conf_matrix = confusion_matrix(y_test, y_pred)
```

0.11 Evaluation Results

```
[85]: print(f'Accuracy: {accuracy:.2f}')
print(f'Confusion Matrix:\n{conf_matrix}')
print('Classification Report:\n',classification_report(y_test, y_pred))
```

```
Accuracy: 0.80
```

```
Confusion Matrix:
```

```
[[929 107]
```

```
 [168 205]]
```

```
Classification Report:
```


	precision	recall	f1-score	support
0	0.85	0.90	0.87	1036
1	0.66	0.55	0.60	373
accuracy			0.80	1409
macro avg	0.75	0.72	0.73	1409

weighted avg 0.80 0.80 0.80 1409

0.12 Confusion Matrix

```
[86]: cm = confusion_matrix(y_test, y_pred, labels=lr_model.classes_)

# format and display the confusion matrix
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=lr_model.
    ↪classes_)
disp.plot(cmap=plt.cm.Blues)
plt.show()
```



The confusion matrix shows true negatives (TN, or a prediction of 929 when the true label is 0), false negatives (FN, prediction=0 true=1) of 168, true positives (TP, prediction=true=1) of 205 and false positives (FP, prediction=1 true=0) of 107. From this, we can get an idea of how the algorithm is performing and compare multiple models.

0.13 Interpretation

True negatives (TN): The model correctly predicted 929 instances where the true label is 0 (no churn). This signifies the number of customers who were correctly identified as not churning. These

are satisfied customers who were retained by the company.

False negatives (FN): The model incorrectly predicted 0 (no churn) when the true label is 1 (churn) in 168 instances. This signifies the number of customers who were incorrectly identified as not churning, leading to missed opportunities for intervention. These are customers who churned despite the model predicting otherwise.

True positives (TP): The model correctly predicted 205 instances where the true label is 1 (churn). This signifies the number of customers who were correctly identified as churning. These are customers who actually churned, and the model successfully flagged them for attention or intervention.

False positives (FP): The model incorrectly predicted 1 (churn) when the true label is 0 (no churn) in 107 instances. This signifies the number of customers who were incorrectly identified as churning, leading to unnecessary intervention or resources being allocated to customers who were not at risk of churning.

0.14 Comparison with No information rate

```
[87]: no_info_rate = max(y_train.value_counts(normalize=True))
      print(f'No Information Rate: {no_info_rate}')
```

No Information Rate: 0.7344692935747249

No-Information Rate (0.734)

This rate serves as a baseline metric, indicating the accuracy achieved by always predicting the majority class (i.e., No churn).

0.15 Fine tuning the model

```
[88]: lr_model.predict_proba(X_test)
```

```
[88]: array([[0.44008218, 0.55991782],
           [0.94083688, 0.05916312],
           [0.99605303, 0.00394697],
           ...,
           [0.94430226, 0.05569774],
           [0.98859175, 0.01140825],
           [0.54155396, 0.45844604]])
```

```
[89]: lr_model.predict(X_test)[:15]
```

```
[89]: array([1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0])
```

he result is an array of binary values (0 or 1), where each value represents the predicted class for a specific instance in the test set. For binary classification problems like predicting churn (1 or 0), 0 represents one class - no churn - and 1 represents the other class - churn.

```
[90]: (lr_model.predict_proba(X_test)[:10, 1] > 0.5).astype('int')
```

```
[90]: array([1, 0, 0, 1, 0, 0, 0, 0, 0, 0])
```

A boolean array indicating whether the predicted probability of class 1 (positive class) for each sample in the first 10 rows of the test set is greater than 0.5. The `.astype('int')` method then converts these boolean values to integers, where True becomes 1 and False becomes 0.

So, the resulting array `[0, 1, 1, 0, 0, 0, 1, 0, 0, 1]` corresponds to the binary predictions (0 or 1) for the first 10 samples in the test set based on the threshold of 0.3. A value of 1 indicates that the predicted probability of belonging to class 1 is greater than 0.3, while a value of 0 indicates that it is not.

0.16 Using least threshold

```
[91]: predictions_lt = (lr_model.predict_proba(X_test)[:10, 1] > 0.15).astype('int')
      predictions_lt
```

```
[91]: array([1, 0, 0, 1, 0, 1, 0, 0, 0, 1])
```

0.17 Accuracy and TP rate

```
[92]: predictions_lt = (lr_model.predict_proba(X_test)[: , 1] > 0.15).astype('int')
      print(accuracy_score(y_test, predictions_lt))
```

```
0.652235628105039
```

The accuracy score of approximately 65.22% signifies the overall proportion of correct predictions made by the model on the test set. Specifically, it indicates that:

Out of all the samples in the test set, approximately 65.22% of them were correctly classified by the model.

The model's predictions matched the true labels for approximately 65.22% of the samples in the test set.

```
[93]: tn, fp, fn, tp = confusion_matrix(y_test, predictions_lt).flatten()
      print(tp / (tp + fn))
```

```
0.9463806970509383
```

The true positive rate (TPR), also known as sensitivity or recall, for a binary classification model is calculated.

`tp` represents the number of true positives (correctly predicted positive instances).

`fn` represents the number of false negatives (actual positive instances incorrectly predicted as negative).

The formula $(tp / (tp + fn))$ calculates the true positive rate, which is the ratio of correctly predicted positive instances to the total actual positive instances.

The true positive rate is approximately 0.9464, indicating that the model correctly predicted around 94.64% of the positive instances in the test set.


```
[94]: lr_model.coef_
```

```
[94]: array([[ -1.80700392,  -0.84090054,   2.47241881,   1.48021476,  -1.79273781,
          -0.06965429,  -0.13987485,  -0.3897276 ,  -0.18019845,  -0.95531267,
           0.01075943,  -0.9660721 ]])
```

```
[95]: coef_df = pd.DataFrame(data=lr_model.coef_, columns=X.columns)

coef_df.T.sort_values(by=0).plot.bar(legend=False)
```

```
[95]: <Axes: >
```



0.18 Interpretation

The plot visualizes the coefficients for each feature in a bar chart format.

Features with positive coefficients are represented by bars pointing upwards, while features with negative coefficients are represented by bars pointing downwards.

The length of each bar indicates the magnitude of the coefficient, representing the strength of the association with the target variable.

Features with longer bars (either positive or negative) such as MonthlyCharges and tenure have a stronger influence on the model's predictions.

Those with positive coefficients - MonthlyCharges and TotalCharges - have a positive impact on the likelihood of churn, while features with negative coefficients (downward bars) have a negative impact.

The plot provides insights into which features are most influential in predicting churn, allowing stakeholders to prioritize and focus on key factors affecting customer churn.

The coefficients and the plot help in understanding the relative importance of different features in predicting churn, guiding decision-making processes aimed at reducing churn rates and improving customer retention strategies.

0.19 Advanced Section

0.20 Comparison with other ML models

```
[99]: rf_model = RandomForestClassifier(random_state=42)
      rf_model.fit(X_train, y_train)
```

```
[99]: RandomForestClassifier(random_state=42)
```

```
[101]: print(rf_model.score(X_train, y_train))
      print(rf_model.score(X_test, y_test))
```

```
0.9941427050053249
0.7849538679914834
```

```
[103]: # Evaluate RFC model
      y_pred_rf = rf_model.predict(X_test)
      accuracy_rf = accuracy_score(y_test, y_pred_rf)

      print(f"Accuracy of KNN: {accuracy_rf}")
```

```
Accuracy of KNN: 0.7849538679914834
```

0.21 Using SVC

```
[100]: svc_model = SVC(kernel='linear', C=1)
       svc_model.fit(X_train, y_train)
```

```
[100]: SVC(C=1, kernel='linear')
```

```
[104]: print(svc_model.score(X_train, y_train))
       print(svc_model.score(X_test, y_test))
```

0.7875399361022364

0.8005677785663591

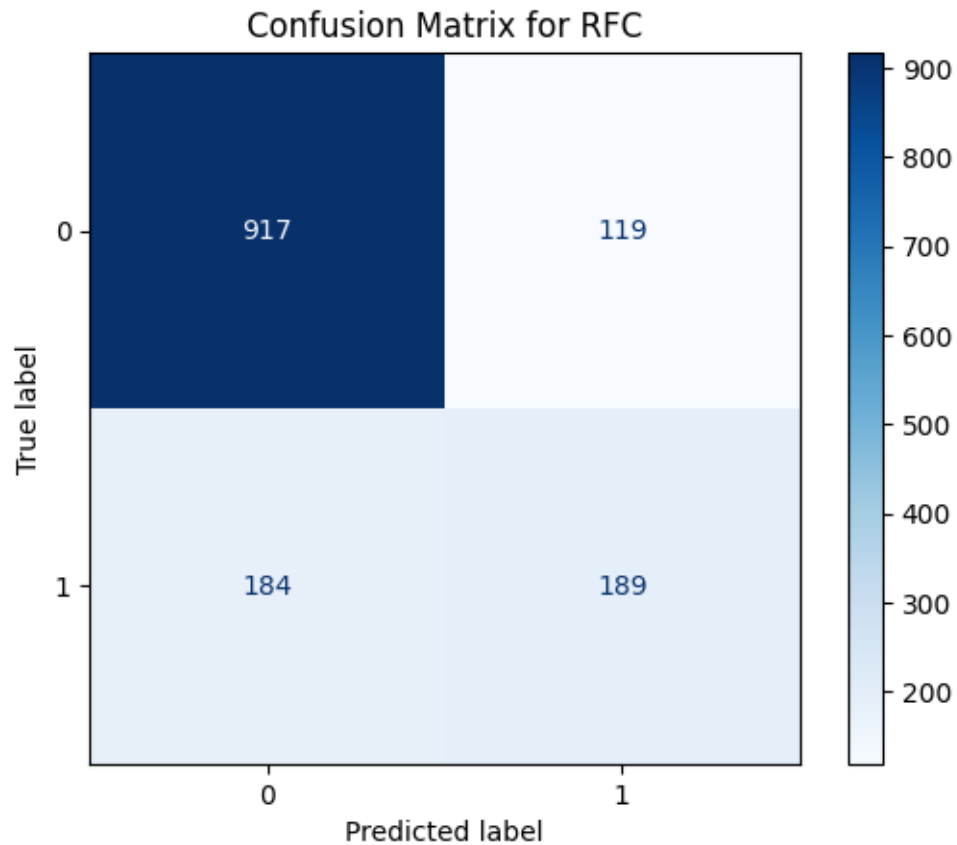
```
[105]: # Evaluate SVC model
       y_pred_svc = svc_model.predict(X_test)
       accuracy_svc = accuracy_score(y_test, y_pred_svc)

       print(f"Accuracy of SVC: {accuracy_svc}")
```

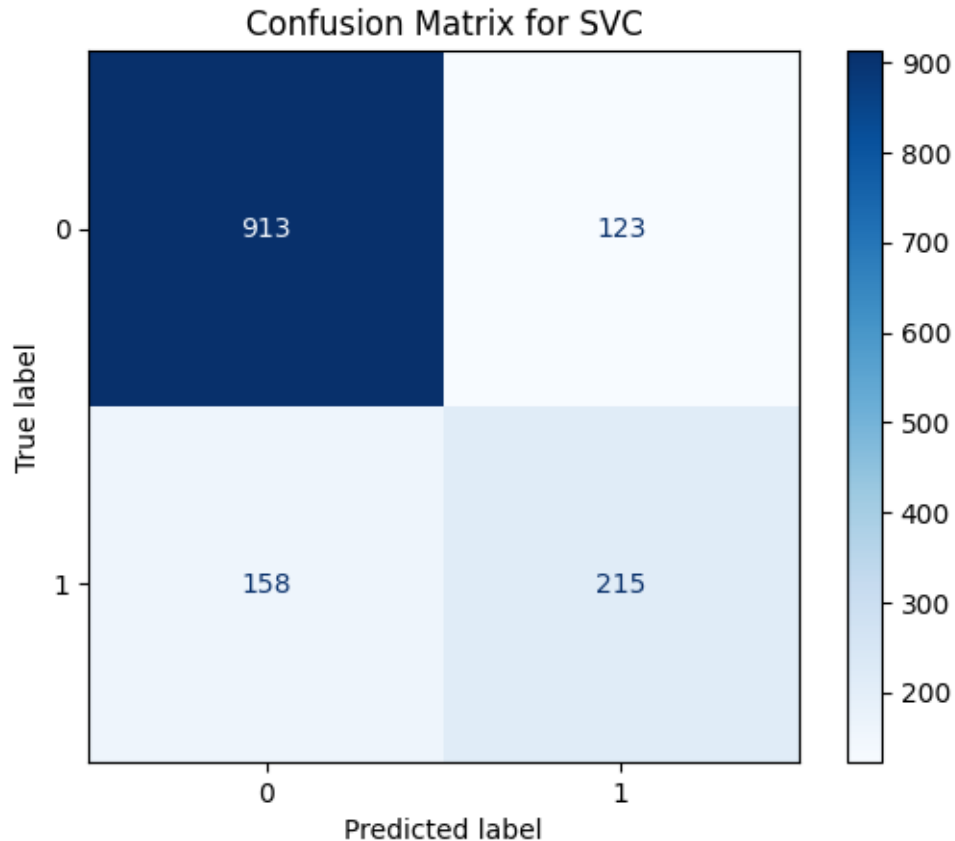
Accuracy of SVC: 0.8005677785663591

0.22 Plot Confusion matrix

```
[107]: # Confusion Matrix for RFC
       cm_rf = confusion_matrix(y_test, y_pred_rf)
       disp_rf = ConfusionMatrixDisplay(confusion_matrix=cm_rf,
       ↪display_labels=rf_model.classes_)
       disp_rf.plot(cmap=plt.cm.Blues)
       plt.title("Confusion Matrix for RFC")
       plt.show()
```



```
[108]: # Confusion Matrix for SVC
cm_svc = confusion_matrix(y_test, y_pred_svc)
disp_svc = ConfusionMatrixDisplay(confusion_matrix=cm_svc,
    ↪display_labels=svc_model.classes_)
disp_svc.plot(cmap=plt.cm.Blues)
plt.title("Confusion Matrix for SVC")
plt.show()
```

0.23 Hyperparameter tuning

```
[109]: # Logistic Regression
lr_params = {'max_iter': [2000, 4000, 6000]}
lr_model = LogisticRegression()
lr_grid = GridSearchCV(lr_model, param_grid=lr_params, cv=3, scoring='accuracy')
lr_grid.fit(X_train, y_train)

# Random Forest Classifier (RFC)
rf_params = {'n_estimators': [50, 100, 200], 'max_depth': [None, 10, 20]}
rf_model = RandomForestClassifier()
rf_grid = GridSearchCV(rf_model, param_grid=rf_params, cv=3, scoring='accuracy')
rf_grid.fit(X_train, y_train)

# Support Vector Classifier (SVC)
svc_params = {'C': [0.1, 1, 10], 'kernel': ['linear', 'rbf']}
svc_model = SVC(probability=True)
svc_grid = GridSearchCV(svc_model, param_grid=svc_params, cv=3,
    ↪scoring='accuracy')
```

```

svc_grid.fit(X_train, y_train)

# Calculate test accuracies
lr_accuracy = accuracy_score(y_test, lr_grid.predict(X_test))
rf_accuracy = accuracy_score(y_test, rf_grid.predict(X_test))
svc_accuracy = accuracy_score(y_test, svc_grid.predict(X_test))

# Display best parameters and accuracy for each model
print("Logistic Regression - Best Parameters:", lr_grid.best_params_)
print("Logistic Regression - Test Accuracy:", lr_accuracy)

print("Random Forest Classifier - Best Parameters:", rf_grid.best_params_)
print("Random Forest Classifier - Test Accuracy:", rf_accuracy)

print("Support Vector Classifier - Best Parameters:", svc_grid.best_params_)
print("Support Vector Classifier - Test Accuracy:", svc_accuracy)

```

```

Logistic Regression - Best Parameters: {'max_iter': 2000}
Logistic Regression - Test Accuracy: 0.8048261178140526
Random Forest Classifier - Best Parameters: {'max_depth': 10, 'n_estimators':
200}
Random Forest Classifier - Test Accuracy: 0.808374733853797
Support Vector Classifier - Best Parameters: {'C': 1, 'kernel': 'linear'}
Support Vector Classifier - Test Accuracy: 0.8005677785663591

```

The results showcase the outcome of hyperparameter tuning and testing for three different machine learning models: Logistic Regression, Random Forest Classifier, and Support Vector Classifier (SVC).

For Logistic Regression, the best parameter configuration was found to be {'max_iter': 2000}, leading to a test accuracy of approximately 80.5%.

The Random Forest Classifier, with optimized parameters {'max_depth': 10, 'n_estimators': 200}, exhibited a test accuracy of around 80.8%.

Lastly, the Support Vector Classifier achieved the best performance with parameters {'C': 1, 'kernel': 'linear'}, resulting in a test accuracy of approximately 80.1%. These findings provide insights into the effectiveness of each model and their respective configurations in making accurate predictions on the test set. The test accuracies are crucial metrics, indicating the models' ability to generalize well to unseen data.

0.24 Determine best model

```

[110]: model_accuracies = {
        'Logistic Regression': lr_accuracy,
        'RFC': rf_accuracy,
        'Support Vector Classifier': svc_accuracy
    }

best_model = max(model_accuracies, key=model_accuracies.get)

```

```
print(f"The best model is: {best_model} with test accuracy_{  
↪{model_accuracies[best_model]:.3f}")
```

The best model is: RFC with test accuracy 0.808

0.25 Summary

Commencing with data cleaning and exploration, we handled non-numeric columns and transformed features to ensure a clean, representative dataset. Feature engineering added depth to the analysis, creating new features and transforming existing ones. After preprocessing the data, we delved into selecting and training various machine learning models, including Random Forest Classifier, LogisticRegression and Support Vector Classifier.

Hyperparameter tuning refined the models, optimizing their performance. Rigorous evaluation on the test set provided a clear picture of model generalization. We explored feature importance through visualizations like bar charts.

Finally, we determined the best-performing model based on test accuracies - which was RFC, completing an approach to building and interpreting a robust machine learning solution.

1 Deployment

1.1 API Integration

The model can be deployed as an API (Application Programming Interface), allowing seamless communication between the customer management system and the predictive model.

Deploying the churn prediction model as an API serves as a strategic move for real-time decision-making within the customer management system. This integration facilitates continuous communication between the predictive model and the customer data, allowing for swift predictions of churn probability scores. By embedding the model into the existing customer management system, the API empowers the system to automatically assess the likelihood of churn for individual customers based on their historical interactions and present behavior.

The integration opens avenues for customer engagement, as customer service representatives can utilize the churn predictions to identify and reach out to customers at risk of leaving. This enables personalized retention efforts, such as tailored offers or issue resolutions, fostering a more responsive and customer-centric approach.

Beyond customer retention, the model's insights also contribute to strategic business decisions. The correlations between specific features or services and customer retention highlighted by the model can inform product development strategies. This, in turn, assists in refining existing offerings or introducing new features that align with customer preferences, ultimately enhancing the overall customer experience.