

Week4_Assignment

February 11, 2024

```
[152]: import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.tree import DecisionTreeClassifier, export_text
from sklearn.ensemble import RandomForestClassifier
import matplotlib.pyplot as plt
import seaborn as sns
```

0.1 Load and preprocess data

```
[153]: df = pd.read_csv("cleaned_churn_data.csv")

df.tail(15)
```

```
[153]:
```

	customerID	tenure	PhoneService	Contract	PaymentMethod \
7028	7029	68	1	3	Bank transfer (automatic)
7029	7030	6	0	0	Electronic check
7030	7031	2	1	0	Mailed check
7031	7032	55	1	1	Credit card (automatic)
7032	7033	1	1	0	Electronic check
7033	7034	38	1	0	Credit card (automatic)
7034	7035	67	1	0	Credit card (automatic)
7035	7036	19	1	0	Bank transfer (automatic)
7036	7037	12	0	1	Electronic check
7037	7038	72	1	3	Bank transfer (automatic)
7038	7039	24	1	1	Mailed check
7039	7040	72	1	1	Credit card (automatic)
7040	7041	11	0	0	Electronic check
7041	7042	4	1	0	Mailed check
7042	7043	66	1	3	Bank transfer (automatic)

	MonthlyCharges	TotalCharges	Churn	MonthlyCharges_to_tenure_Ratio
7028	64.10	4326.25	0	0.942647
7029	44.40	263.05	0	7.400000
7030	20.05	39.25	0	10.025000
7031	60.00	3316.10	0	1.090909
7032	75.75	75.75	1	75.750000

7033	69.50	2625.25	0	1.828947
7034	102.95	6886.25	1	1.536567
7035	78.70	1495.10	0	4.142105
7036	60.65	743.30	0	5.054167
7037	21.15	1419.40	0	0.293750
7038	84.80	1990.50	0	3.533333
7039	103.20	7362.90	0	1.433333
7040	29.60	346.45	0	2.690909
7041	74.40	306.60	1	18.600000
7042	105.65	6844.50	0	1.600758

```
[154]: # Drop customerID - not useful for modeling
df = df.drop('customerID', axis=1)
df
```

```
[154]:
```

	tenure	PhoneService	Contract	PaymentMethod \
0	1	0	0	Electronic check
1	34	1	1	Mailed check
2	2	1	0	Mailed check
3	45	0	1	Bank transfer (automatic)
4	2	1	0	Electronic check
...
7038	24	1	1	Mailed check
7039	72	1	1	Credit card (automatic)
7040	11	0	0	Electronic check
7041	4	1	0	Mailed check
7042	66	1	3	Bank transfer (automatic)

	MonthlyCharges	TotalCharges	Churn	MonthlyCharges_to_tenure_Ratio
0	29.85	29.85	0	29.850000
1	56.95	1889.50	0	1.675000
2	53.85	108.15	1	26.925000
3	42.30	1840.75	0	0.940000
4	70.70	151.65	1	35.350000
...
7038	84.80	1990.50	0	3.533333
7039	103.20	7362.90	0	1.433333
7040	29.60	346.45	0	2.690909
7041	74.40	306.60	1	18.600000
7042	105.65	6844.50	0	1.600758

[7043 rows x 8 columns]

```
[155]: payment_method_dummies = pd.get_dummies(df['PaymentMethod'])
df = pd.concat([df, payment_method_dummies], axis=1)
df.head()
```

```
[155]:
```

	tenure	PhoneService	Contract	PaymentMethod	MonthlyCharges \
0	1	0	0	Electronic check	29.85
1	34	1	1	Mailed check	56.95
2	2	1	0	Mailed check	53.85
3	45	0	1	Bank transfer (automatic)	42.30
4	2	1	0	Electronic check	70.70

	TotalCharges	Churn	MonthlyCharges_to_tenure_Ratio \
0	29.85	0	29.850
1	1889.50	0	1.675
2	108.15	1	26.925
3	1840.75	0	0.940
4	151.65	1	35.350

	Bank transfer (automatic)	Credit card (automatic)	Electronic check \
0	False	False	True
1	False	False	False
2	False	False	False
3	True	False	False
4	False	False	True

	Mailed check
0	False
1	True
2	True
3	False
4	False

```
[156]: dummies = ['Electronic check', 'Mailed check', 'Bank transfer (automatic)',
↳ 'Credit card (automatic)']

for column in dummies:
    df[column] = pd.factorize(df[column])[0]

df.sample(5)
```

```
[156]:
```

	tenure	PhoneService	Contract	PaymentMethod \
2088	36	1	1	Electronic check
153	62	1	3	Electronic check
6170	1	1	0	Credit card (automatic)
2178	53	1	3	Credit card (automatic)
619	7	1	0	Bank transfer (automatic)

	MonthlyCharges	TotalCharges	Churn	MonthlyCharges_to_tenure_Ratio \
2088	55.80	1941.50	0	1.550000
153	86.10	5215.25	0	1.388710
6170	19.40	19.40	0	19.400000

2178	19.90	1110.05	0	0.375472
619	78.55	522.95	0	11.221429

	Bank transfer (automatic)	Credit card (automatic)	Electronic check	\
2088	0	0	0	
153	0	0	0	
6170	0	1	1	
2178	0	1	1	
619	1	0	1	

	Mailed check
2088	0
153	0
6170	0
2178	0
619	0

```
[157]: df = df.drop('PaymentMethod', axis=1)
df
```

```
[157]:
```

	tenure	PhoneService	Contract	MonthlyCharges	TotalCharges	Churn	\
0	1	0	0	29.85	29.85	0	
1	34	1	1	56.95	1889.50	0	
2	2	1	0	53.85	108.15	1	
3	45	0	1	42.30	1840.75	0	
4	2	1	0	70.70	151.65	1	
...	
7038	24	1	1	84.80	1990.50	0	
7039	72	1	1	103.20	7362.90	0	
7040	11	0	0	29.60	346.45	0	
7041	4	1	0	74.40	306.60	1	
7042	66	1	3	105.65	6844.50	0	

	MonthlyCharges_to_tenure_Ratio	Bank transfer (automatic)	\
0	29.850000	0	
1	1.675000	0	
2	26.925000	0	
3	0.940000	1	
4	35.350000	0	
...	
7038	3.533333	0	
7039	1.433333	0	
7040	2.690909	0	
7041	18.600000	0	
7042	1.600758	1	

	Credit card (automatic)	Electronic check	Mailed check
--	-------------------------	------------------	--------------

0	0	0	0
1	0	1	1
2	0	1	1
3	0	1	0
4	0	0	0
...
7038	0	1	1
7039	1	1	0
7040	0	0	0
7041	0	1	1
7042	0	1	0

[7043 rows x 11 columns]

```
[158]: df.isna().sum()
```

```
[158]: tenure                0
PhoneService                0
Contract                    0
MonthlyCharges              0
TotalCharges                0
Churn                       0
MonthlyCharges_to_tenure_Ratio  0
Bank transfer (automatic)    0
Credit card (automatic)     0
Electronic check             0
Mailed check                 0
dtype: int64
```

0.2 Remove outliers

```
[159]: numerical_columns = df.select_dtypes(include=[np.number]).columns

# Function to remove outliers based on IQR
def remove_outliers_iqr(data_frame, columns):
    for column in columns:
        Q1 = data_frame[column].quantile(0.25)
        Q3 = data_frame[column].quantile(0.75)
        IQR = Q3 - Q1
        lower_bound = Q1 - 1.5 * IQR
        upper_bound = Q3 + 1.5 * IQR
        data_frame = data_frame[(data_frame[column] >= lower_bound) &
↪ (data_frame[column] <= upper_bound)]
    return data_frame

df_clean = remove_outliers_iqr(df, numerical_columns)
df_clean
```

```
[159]:
```

	tenure	PhoneService	Contract	MonthlyCharges	TotalCharges	Churn	\
1	34	1	1	56.95	1889.50	0	
5	8	1	0	99.65	820.50	1	
8	28	1	0	104.80	3046.05	1	
10	13	1	0	49.95	587.45	0	
14	25	1	0	105.50	2686.05	0	
...	
7023	63	1	0	103.50	6479.40	0	
7027	13	1	0	73.35	931.55	0	
7030	2	1	0	20.05	39.25	0	
7038	24	1	1	84.80	1990.50	0	
7041	4	1	0	74.40	306.60	1	

	MonthlyCharges_to_tenure_Ratio	Bank transfer (automatic)	\
1	1.675000	0	
5	12.456250	0	
8	3.742857	0	
10	3.842308	0	
14	4.220000	0	
...	
7023	1.642857	0	
7027	5.642308	0	
7030	10.025000	0	
7038	3.533333	0	
7041	18.600000	0	

	Credit card (automatic)	Electronic check	Mailed check
1	0	1	1
5	0	0	0
8	0	0	0
10	0	1	1
14	0	0	0
...
7023	0	0	0
7027	0	1	1
7030	0	1	1
7038	0	1	1
7041	0	1	1


```
[2523 rows x 11 columns]
```

0.3 Split data into features and targets

```
[160]: X = df.drop('Churn', axis=1)
y = df['Churn']
```

0.4 Split data into training and test sets

```
[161]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,  
↳ random_state=42)
```

0.5 Fix infinity values error in the dataset

```
[162]: print("Infinity values in X_train:", np.any(np.isinf(X_train)))  
print("NaN values in X_train:", np.any(np.isnan(X_train)))
```

Infinity values in X_train: True

NaN values in X_train: False

0.6 Handle infinity values

```
[163]: columns_with_infinity = X_train.columns[np.isinf(X_train).any()]  
  
print("Columns with Infinity values:", columns_with_infinity)  
  
# Replace infinity values  
X_train[columns_with_infinity] = X_train[columns_with_infinity].replace([np.  
↳ inf, -np.inf], np.nan)
```

Columns with Infinity values: Index(['MonthlyCharges_to_tenure_Ratio'],
dtype='object')

0.7 Align the data to avoid mismatch

```
[164]: print("Shape of X_train:", X_train.shape)  
print("Shape of y_train:", y_train.shape)  
  
print("Duplicate index values in X_train:", X_train.index.duplicated().any())  
print("Duplicate index values in y_train:", y_train.index.duplicated().any())  
  
# Reindex y_train to match X_train  
y_train = y_train.reindex(X_train.index)
```

Shape of X_train: (5634, 10)

Shape of y_train: (5634,)

Duplicate index values in X_train: False

Duplicate index values in y_train: False

```
[165]: print("Shape of X_train:", X_train.shape)  
print("Shape of y_train:", y_train.shape)
```

Shape of X_train: (5634, 10)

Shape of y_train: (5634,)

0.8 Fit and plot DT

```
[166]: dt_model = DecisionTreeClassifier(max_depth=3)
dt_model.fit(X_train, y_train)
tree_rules = export_text(dt_model, feature_names=list(X.columns))
print(tree_rules)
```

```
|--- Contract <= 0.50
|   |--- MonthlyCharges_to_tenure_Ratio <= 7.84
|   |   |--- MonthlyCharges <= 69.97
|   |   |   |--- class: 0
|   |   |--- MonthlyCharges > 69.97
|   |   |   |--- class: 0
|   |--- MonthlyCharges_to_tenure_Ratio > 7.84
|   |   |--- MonthlyCharges <= 67.30
|   |   |   |--- class: 0
|   |   |--- MonthlyCharges > 67.30
|   |   |   |--- class: 1
|--- Contract > 0.50
|   |--- MonthlyCharges <= 93.67
|   |   |--- Contract <= 2.00
|   |   |   |--- class: 0
|   |   |--- Contract > 2.00
|   |   |   |--- class: 0
|   |--- MonthlyCharges > 93.67
|   |   |--- Contract <= 2.00
|   |   |   |--- class: 0
|   |   |--- Contract > 2.00
|   |   |   |--- class: 0
```

First Decision Node (Contract <= 0.50)

If the contract duration is short (month-to-month contracts)

The model looks at the ratio of MonthlyCharges to tenure. If this ratio is low, and MonthlyCharges is low, predict class 0 (indicating no churn).

If this ratio is high, and MonthlyCharges is high, predict class 0 (indicating no churn).

If this ratio is high, and MonthlyCharges is relatively moderate, predict class 1 (indicating potential churn).

Second Decision Node (Contract > 0.50)

If the contract duration is longer (one or two-year contracts)

The model looks at the MonthlyCharges. If MonthlyCharges is low, predict class 0 (no churn).

If MonthlyCharges is high, the model further considers the contract duration.

If the contract duration is short (less than or equal to 2.00), predict class 0 (no churn).

If the contract duration is long (greater than 2.00), predict class 0 (no churn).

0.9 Interpretation

Short-term or month-to-month contracts with high monthly charges and a moderate ratio of monthly charges to tenure are associated with a higher likelihood of churn.

Longer-term contracts with low or high monthly charges are less likely to result in churn, regardless of the contract duration.

These interpretations provide insights into the factors that the model considers important for predicting customer churn, and they align with the business understanding of wanting to identify potential churners.

0.10 Hyperparameter tuning for DT

```
[167]: param_grid = {'max_depth': [3, 5, 7, 10]}
dt_model = DecisionTreeClassifier()
grid_search = GridSearchCV(dt_model, param_grid, cv=5)
grid_search.fit(X_train, y_train)

best_max_depth = grid_search.best_params_['max_depth']

print(best_max_depth)
```

3

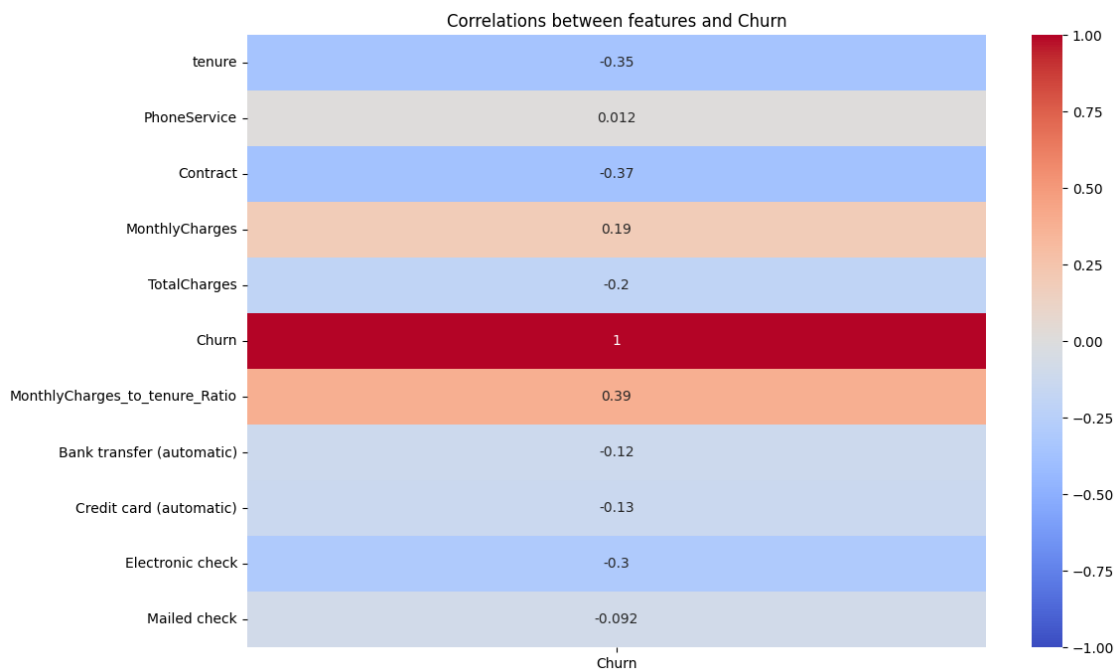
```
[168]: dt_model_tuned = DecisionTreeClassifier(max_depth=best_max_depth)
dt_model_tuned.fit(X_train, y_train)
tree_rules_tuned = export_text(dt_model_tuned, feature_names=list(X.columns))
print(tree_rules_tuned)
```

```
|--- Contract <= 0.50
|   |--- MonthlyCharges_to_tenure_Ratio <= 7.84
|   |   |--- MonthlyCharges <= 69.97
|   |   |   |--- class: 0
|   |   |--- MonthlyCharges > 69.97
|   |   |   |--- class: 0
|   |--- MonthlyCharges_to_tenure_Ratio > 7.84
|   |   |--- MonthlyCharges <= 67.30
|   |   |   |--- class: 0
|   |   |--- MonthlyCharges > 67.30
|   |   |   |--- class: 1
|--- Contract > 0.50
|   |--- MonthlyCharges <= 93.67
|   |   |--- Contract <= 2.00
|   |   |   |--- class: 0
|   |   |--- Contract > 2.00
|   |   |   |--- class: 0
|   |--- MonthlyCharges > 93.67
|   |   |--- Contract <= 2.00
|   |   |   |--- class: 0
```

```
|    |    |--- Contract > 2.00
|    |    |--- class: 0
```

0.11 Correlations between features and targets

```
[169]: plt.figure(figsize=(12, 8))
sns.heatmap(df.corr()[['Churn']], annot=True, cmap='coolwarm', vmin=-1, vmax=1)
plt.title("Correlations between features and Churn")
plt.show()
```



0.12 Handle Nan and infinity values in X_train

```
[170]: nan_values_in_X_train = X_train.isna().any().any()
print("NaN values in X_train:", nan_values_in_X_train)

# Check for Infinite values in X_train
infinite_values_in_X_train = np.isfinite(X_train).all().all()
print("Infinity values in X_train:", not infinite_values_in_X_train)

# Check for Constant features in X_train
constant_features_in_X_train = X_train.columns[X_train.nunique() == 1].tolist()
print("Constant features in X_train:", constant_features_in_X_train)
```

NaN values in X_train: True

Infinity values in X_train: True

Constant features in X_train: []

0.13 Fill nan values with mean

```
[171]: X_train_filled = X_train.fillna(X_train.mean())

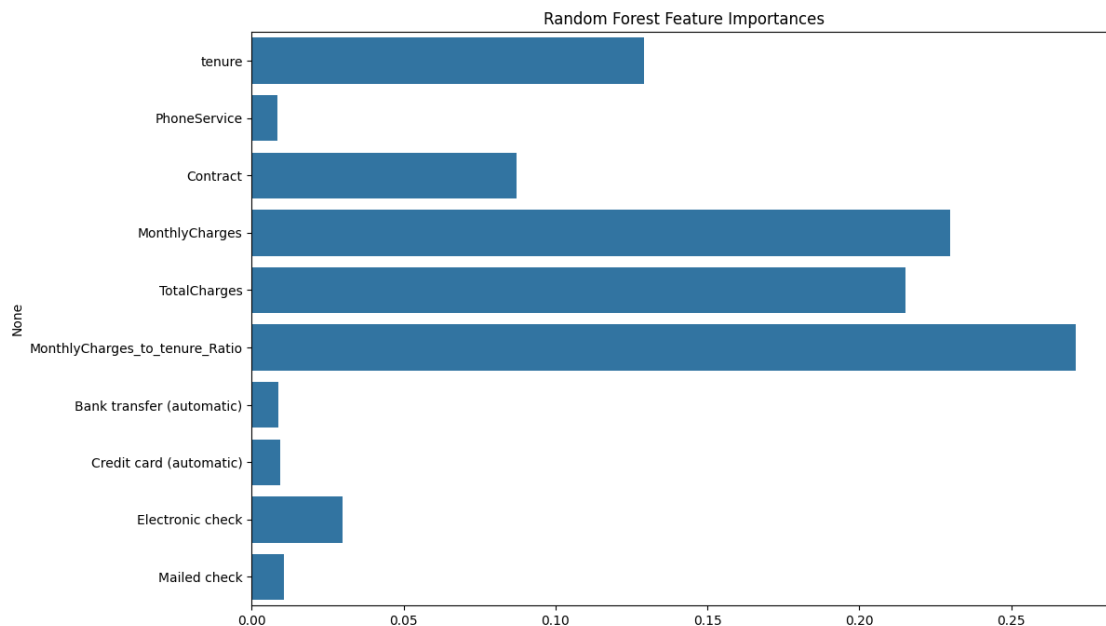
# Replace infinity values with a large finite value
X_train_filled.replace([np.inf, -np.inf], np.finfo(np.float64).max,
    ↪ inplace=True)
```

0.14 Fit and plot model

```
[173]: rf_model = RandomForestClassifier(n_estimators=100, random_state=42)
rf_model.fit(X_train_filled, y_train)
```

```
[173]: RandomForestClassifier(random_state=42)
```

```
[174]: plt.figure(figsize=(12, 8))
sns.barplot(x=rf_model.feature_importances_, y=X.columns)
plt.title("Random Forest Feature Importances")
plt.show()
```



0.15 Remove less important features

```
[175]: less_important_features = ['PhoneService', 'Bank transfer (automatic)', 'Credit_
↳ card (automatic)', 'Mailed check']

X_train_filtered = X_train_filled.drop(less_important_features, axis=1)
X_test_filtered = X_test.drop(less_important_features, axis=1)
X_train_filtered.shape
```

```
[175]: (5634, 6)
```

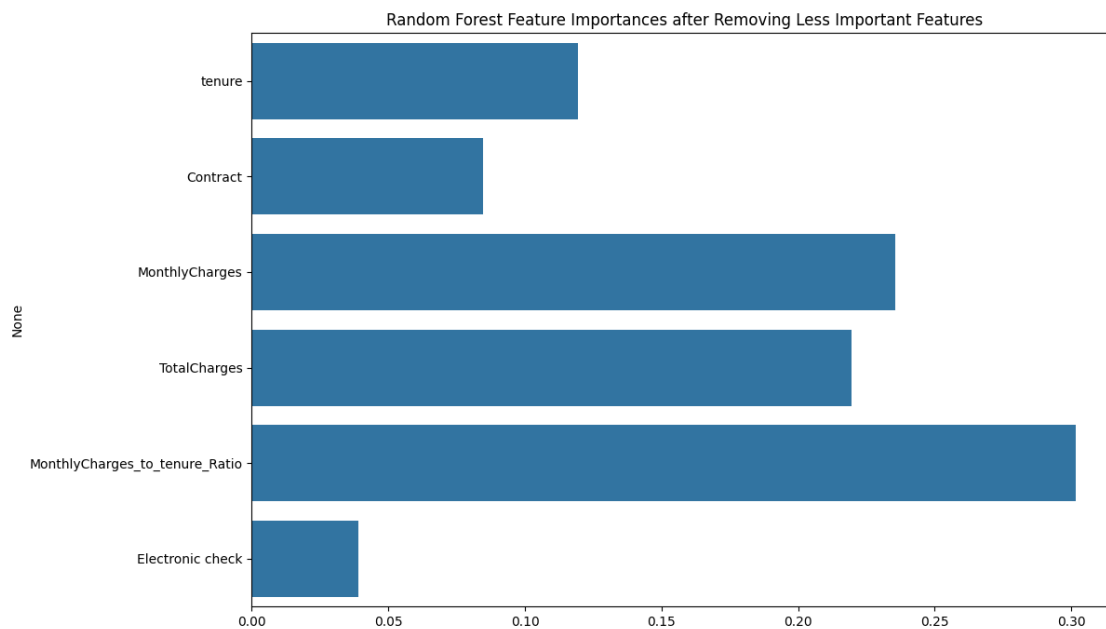
0.16 Fit to new data

```
[176]: rf_model_filtered = RandomForestClassifier(n_estimators=100, random_state=42)
rf_model_filtered.fit(X_train_filtered, y_train)
```

```
[176]: RandomForestClassifier(random_state=42)
```

0.17 Plot feature importances

```
[177]: plt.figure(figsize=(12, 8))
sns.barplot(x=rf_model_filtered.feature_importances_, y=X_train_filtered.
↳ columns)
plt.title("Random Forest Feature Importances after Removing Less Important_
↳ Features")
plt.show()
```



0.18 Summary

We perform a comprehensive analysis on customer churn data, focusing on building, tuning, and evaluating machine learning models. Initial data preprocessing involves loading, dummy encoding, and factorization of categorical features. We then ensure data integrity by handling Infinity and NaN values.

We then construct an initial decision tree model, conduct hyperparameter tuning using GridSearchCV, and visualize the decision tree rules. Additionally, we plot correlations between features and the target variable (Churn). Random Forest modeling is introduced, featuring an initial model with visualized feature importances. We then identify and remove less-important features, leading to a new Random Forest model. The overall goal is to enhance the predictive capabilities of these models for customer churn prediction in a telecommunications company.