

Week4_Assignment

February 12, 2024

```
[1]: import pandas as pd
import h2o
from h2o.estimators import H2ORandomForestEstimator
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.tree import DecisionTreeClassifier, export_text
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
```

0.1 Load data

```
[2]: df = pd.read_csv("preped_churn_data.csv")

df.sample(10)
```

```
[2]:      tenure PhoneService      Contract      PaymentMethod \
1198      42         Yes      Two year  Bank transfer (automatic)
5028      72         Yes      Two year      Electronic check
1532      14         Yes  Month-to-month  Credit card (automatic)
83        42         Yes  Month-to-month  Bank transfer (automatic)
4003      38         Yes      Two year      Mailed check
1681      72         Yes      Two year  Credit card (automatic)
879       33         Yes  Month-to-month  Bank transfer (automatic)
2921       9         Yes  Month-to-month      Electronic check
5185      50         Yes      One year  Credit card (automatic)
1257      45         No      One year      Electronic check
```

```
      MonthlyCharges  TotalCharges  Churn \
1198           20.75         844.45   No
5028          109.90        7624.20   No
1532           19.60         300.40   No
83            103.80        4327.50   No
4003           20.30         743.05   No
1681           25.00        1849.20   No
879            54.65        1665.20   No
2921           95.50         829.10  Yes
5185           70.80        3478.15   No
```

1257	34.20	1596.60	No
------	-------	---------	----

	TotalCharges_to_MonthlyCharges_ratio	customerID
1198	40.696386	1199
5028	69.373976	5029
1532	15.326531	1533
83	41.690751	84
4003	36.603448	4004
1681	73.968000	1682
879	30.470265	880
2921	8.681675	2922
5185	49.126412	5186
1257	46.684211	1258

```
[3]: df['Churn'] = df['Churn'].map({'Yes': 1, 'No': 0})

df['PhoneService'] = df['PhoneService'].map({'Yes': 1, 'No': 0})
df
```

```
[3]:
```

	tenure	PhoneService	Contract	PaymentMethod \
0	1	0	Month-to-month	Electronic check
1	34	1	One year	Mailed check
2	2	1	Month-to-month	Mailed check
3	45	0	One year	Bank transfer (automatic)
4	2	1	Month-to-month	Electronic check
...
7038	24	1	One year	Mailed check
7039	72	1	One year	Credit card (automatic)
7040	11	0	Month-to-month	Electronic check
7041	4	1	Month-to-month	Mailed check
7042	66	1	Two year	Bank transfer (automatic)

	MonthlyCharges	TotalCharges	Churn \
0	29.85	29.85	0
1	56.95	1889.50	0
2	53.85	108.15	1
3	42.30	1840.75	0
4	70.70	151.65	1
...
7038	84.80	1990.50	0
7039	103.20	7362.90	0
7040	29.60	346.45	0
7041	74.40	306.60	1
7042	105.65	6844.50	0

	TotalCharges_to_MonthlyCharges_ratio	customerID
0	1.000000	1

1	33.178227	2
2	2.008357	3
3	43.516548	4
4	2.144979	5
...
7038	23.472877	7039
7039	71.345930	7040
7040	11.704392	7041
7041	4.120968	7042
7042	64.784666	7043

[7043 rows x 9 columns]

```
[4]: PM_dummies = pd.get_dummies(df['PaymentMethod'], prefix='PaymentMethod')
C_dummies = pd.get_dummies(df['Contract'], prefix='Contract')

df = pd.concat([df, PM_dummies, C_dummies], axis=1)

df = df.loc[:, ~df.columns.duplicated()]

df.head(5)
```

```
[4]:
```

	tenure	PhoneService	Contract	PaymentMethod	\
0	1	0	Month-to-month	Electronic check	
1	34	1	One year	Mailed check	
2	2	1	Month-to-month	Mailed check	
3	45	0	One year	Bank transfer (automatic)	
4	2	1	Month-to-month	Electronic check	

	MonthlyCharges	TotalCharges	Churn	TotalCharges_to_MonthlyCharges_ratio	\
0	29.85	29.85	0	1.000000	
1	56.95	1889.50	0	33.178227	
2	53.85	108.15	1	2.008357	
3	42.30	1840.75	0	43.516548	
4	70.70	151.65	1	2.144979	

	customerID	PaymentMethod_Bank transfer (automatic)	\
0	1	False	
1	2	False	
2	3	False	
3	4	True	
4	5	False	

	PaymentMethod_Credit card (automatic)	PaymentMethod_Electronic check	\
0	False	True	
1	False	False	
2	False	False	

3		False	False
4		False	True

	PaymentMethod_Mailed check	Contract_Month-to-month	Contract_One year \
0	False	True	False
1	True	False	True
2	True	True	False
3	False	False	True
4	False	True	False

	Contract_Two year
0	False
1	False
2	False
3	False
4	False

```
[5]: dummies = ['Contract_Month-to-month', 'Contract_One year', 'Contract_Two year',
↳ 'PaymentMethod_Bank transfer (automatic)', 'PaymentMethod_Electronic check',
↳ 'PaymentMethod_Mailed check', 'PaymentMethod_Credit card (automatic)']

for column in dummies:
    df[column] = pd.factorize(df[column])[0]

df.sample(5)
```

```
[5]:
```

	tenure	PhoneService	Contract	PaymentMethod \
5246	2	1	Month-to-month	Electronic check
3896	38	1	One year	Mailed check
5493	6	1	Month-to-month	Electronic check
2498	27	1	Month-to-month	Electronic check
3552	68	1	Two year	Bank transfer (automatic)

	MonthlyCharges	TotalCharges	Churn \
5246	79.75	164.50	1
3896	20.30	749.35	0
5493	84.35	474.90	1
2498	53.80	1389.85	0
3552	102.10	7149.35	0

	TotalCharges_to_MonthlyCharges_ratio	customerID \
5246	2.062696	5247
3896	36.913793	3897
5493	5.630113	5494
2498	25.833643	2499
3552	70.023017	3553

	PaymentMethod_Bank transfer (automatic) \
5246	0
3896	0
5493	0
2498	0
3552	1

	PaymentMethod_Credit card (automatic)	PaymentMethod_Electronic check \
5246	0	0
3896	0	1
5493	0	0
2498	0	0
3552	0	1

	PaymentMethod_Mailed check	Contract_Month-to-month	Contract_One year \
5246	0	0	0
3896	1	1	1
5493	0	0	0
2498	0	0	0
3552	0	1	0

	Contract_Two year
5246	0
3896	0
5493	0
2498	0
3552	1

```
[6]: df = df.drop(['PaymentMethod', 'Contract', 'customerID'], axis=1)
df
```

```
[6]:
```

	tenure	PhoneService	MonthlyCharges	TotalCharges	Churn \
0	1	0	29.85	29.85	0
1	34	1	56.95	1889.50	0
2	2	1	53.85	108.15	1
3	45	0	42.30	1840.75	0
4	2	1	70.70	151.65	1
...
7038	24	1	84.80	1990.50	0
7039	72	1	103.20	7362.90	0
7040	11	0	29.60	346.45	0
7041	4	1	74.40	306.60	1
7042	66	1	105.65	6844.50	0

	TotalCharges_to_MonthlyCharges_ratio \
0	1.000000
1	33.178227

2	2.008357
3	43.516548
4	2.144979
...	...
7038	23.472877
7039	71.345930
7040	11.704392
7041	4.120968
7042	64.784666

	PaymentMethod_Bank transfer (automatic) \
0	0
1	0
2	0
3	1
4	0
...	...
7038	0
7039	0
7040	0
7041	0
7042	1

	PaymentMethod_Credit card (automatic)	PaymentMethod_Electronic check \
0	0	0
1	0	1
2	0	1
3	0	1
4	0	0
...
7038	0	1
7039	1	1
7040	0	0
7041	0	1
7042	0	1

	PaymentMethod_Mailed check	Contract_Month-to-month	Contract_One year \
0	0	0	0
1	1	1	1
2	1	0	0
3	0	1	1
4	0	0	0
...
7038	1	1	1
7039	0	1	1
7040	0	0	0
7041	1	0	0

7042	0	1	0
------	---	---	---

	Contract_Two year
0	0
1	0
2	0
3	0
4	0
...	...
7038	0
7039	0
7040	0
7041	0
7042	1

[7043 rows x 13 columns]

```
[7]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 13 columns):
#   Column                                          Non-Null Count  Dtype
---  -
0   tenure                                         7043 non-null   int64
1   PhoneService                                  7043 non-null   int64
2   MonthlyCharges                               7043 non-null   float64
3   TotalCharges                                 7043 non-null   float64
4   Churn                                          7043 non-null   int64
5   TotalCharges_to_MonthlyCharges_ratio         7043 non-null   float64
6   PaymentMethod_Bank transfer (automatic)      7043 non-null   int64
7   PaymentMethod_Credit card (automatic)        7043 non-null   int64
8   PaymentMethod_Electronic check               7043 non-null   int64
9   PaymentMethod_Mailed check                   7043 non-null   int64
10  Contract_Month-to-month                      7043 non-null   int64
11  Contract_One year                            7043 non-null   int64
12  Contract_Two year                            7043 non-null   int64
dtypes: float64(3), int64(10)
memory usage: 715.4 KB
```

```
[8]: df.isna().sum()
```

```
[8]: tenure                0
PhoneService              0
MonthlyCharges            0
TotalCharges              0
Churn                     0
TotalCharges_to_MonthlyCharges_ratio  0
```

```

PaymentMethod_Bank transfer (automatic)    0
PaymentMethod_Credit card (automatic)      0
PaymentMethod_Electronic check             0
PaymentMethod_Mailed check                 0
Contract_Month-to-month                    0
Contract_One year                           0
Contract_Two year                           0
dtype: int64

```

0.2 Split into features and targets

```

[9]: X = df.drop('Churn', axis=1)
     y = df['Churn']

```

0.3 Split into training and test sets

```

[10]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    ↪random_state=42)

```

0.4 Fit and plot the DT

```

[11]: dt_model = DecisionTreeClassifier(max_depth=5)
     dt_model.fit(X_train, y_train)

```

```

[11]: DecisionTreeClassifier(max_depth=5)

```

```

[12]: tr= export_text(dt_model, feature_names=list(X.columns))
     print(tr)

```

```

|--- Contract_Month-to-month <= 0.50
|   |--- MonthlyCharges <= 68.62
|   |   |--- TotalCharges_to_MonthlyCharges_ratio <= 3.75
|   |   |   |--- MonthlyCharges <= 20.88
|   |   |   |   |--- TotalCharges <= 20.57
|   |   |   |   |   |--- class: 0
|   |   |   |   |   |--- TotalCharges > 20.57
|   |   |   |   |   |--- class: 0
|   |   |   |--- MonthlyCharges > 20.88
|   |   |   |--- PaymentMethod_Electronic check <= 0.50
|   |   |   |   |--- class: 1
|   |   |   |   |--- PaymentMethod_Electronic check > 0.50
|   |   |   |   |--- class: 0
|   |   |--- TotalCharges_to_MonthlyCharges_ratio > 3.75
|   |   |--- PhoneService <= 0.50
|   |   |   |--- TotalCharges <= 330.30
|   |   |   |   |--- class: 0
|   |   |   |   |--- TotalCharges > 330.30
|   |   |   |   |--- class: 0

```



```

| | | |--- PhoneService > 0.50
| | | |--- TotalCharges_to_MonthlyCharges_ratio <= 17.79
| | | |--- class: 0
| | | |--- TotalCharges_to_MonthlyCharges_ratio > 17.79
| | | |--- class: 0
| |--- MonthlyCharges > 68.62
| | |--- TotalCharges_to_MonthlyCharges_ratio <= 13.16
| | |--- TotalCharges <= 120.00
| | |--- MonthlyCharges <= 69.88
| | | |--- class: 1
| | | |--- MonthlyCharges > 69.88
| | | |--- class: 1
| | |--- TotalCharges > 120.00
| | |--- MonthlyCharges <= 90.78
| | | |--- class: 1
| | |--- MonthlyCharges > 90.78
| | | |--- class: 1
| | |--- TotalCharges_to_MonthlyCharges_ratio > 13.16
| | |--- PaymentMethod_Electronic check <= 0.50
| | |--- tenure <= 55.50
| | | |--- class: 1
| | |--- tenure > 55.50
| | | |--- class: 0
| | |--- PaymentMethod_Electronic check > 0.50
| | |--- TotalCharges_to_MonthlyCharges_ratio <= 53.08
| | | |--- class: 0
| | |--- TotalCharges_to_MonthlyCharges_ratio > 53.08
| | | |--- class: 0
|--- Contract_Month-to-month > 0.50
| |--- MonthlyCharges <= 93.67
| | |--- Contract_One year <= 0.50
| | |--- MonthlyCharges <= 79.72
| | |--- MonthlyCharges <= 79.67
| | | |--- class: 0
| | |--- MonthlyCharges > 79.67
| | | |--- class: 0
| | |--- MonthlyCharges > 79.72
| | | |--- class: 0
| | |--- Contract_One year > 0.50
| | |--- MonthlyCharges <= 41.38
| | |--- TotalCharges <= 37.08
| | | |--- class: 0
| | |--- TotalCharges > 37.08
| | | |--- class: 0
| | |--- MonthlyCharges > 41.38
| | |--- MonthlyCharges <= 42.70
| | | |--- class: 1
| | |--- MonthlyCharges > 42.70

```

```

|   |   |   |   |   |--- class: 0
|   |--- MonthlyCharges > 93.67
|   |   |--- TotalCharges_to_MonthlyCharges_ratio <= 70.35
|   |   |   |--- Contract_One year <= 0.50
|   |   |   |   |--- MonthlyCharges <= 94.22
|   |   |   |   |   |--- class: 1
|   |   |   |   |   |--- MonthlyCharges > 94.22
|   |   |   |   |   |   |--- class: 0
|   |   |   |   |--- Contract_One year > 0.50
|   |   |   |   |--- TotalCharges <= 7523.15
|   |   |   |   |   |--- class: 0
|   |   |   |   |   |--- TotalCharges > 7523.15
|   |   |   |   |   |   |--- class: 1
|   |   |--- TotalCharges_to_MonthlyCharges_ratio > 70.35
|   |   |   |--- TotalCharges <= 8639.60
|   |   |   |   |--- TotalCharges_to_MonthlyCharges_ratio <= 74.17
|   |   |   |   |   |--- class: 0
|   |   |   |   |   |--- TotalCharges_to_MonthlyCharges_ratio > 74.17
|   |   |   |   |   |   |--- class: 0
|   |   |   |--- TotalCharges > 8639.60
|   |   |   |   |--- class: 1

```

0.5 Hyperparameter tuning

```

[13]: param_grid = {'max_depth': [3, 5, 7, 10]}
      dt_model = DecisionTreeClassifier()
      grid_search = GridSearchCV(dt_model, param_grid, cv=5)
      grid_search.fit(X_train, y_train)

```

```

[13]: GridSearchCV(cv=5, estimator=DecisionTreeClassifier(),
                  param_grid={'max_depth': [3, 5, 7, 10]})

```

0.6 Bests max_depth

```

[14]: best_max_depth = grid_search.best_params_['max_depth']
      best_max_depth

```

```

[14]: 3

```

0.7 Fit and plot DT with best hyperparameters

```

[15]: dt_model_tuned = DecisionTreeClassifier(max_depth=best_max_depth)
      dt_model_tuned.fit(X_train, y_train)

```

```

[15]: DecisionTreeClassifier(max_depth=3)

```

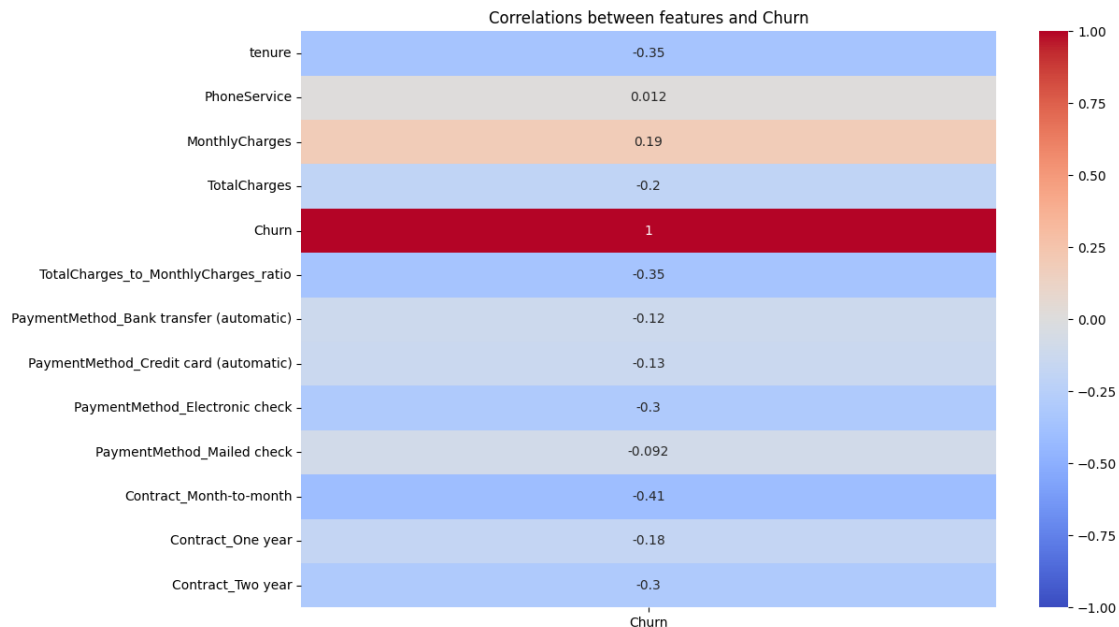
```
[16]: tr_tuned = export_text(dt_model_tuned, feature_names=list(X.columns))
      print(tr_tuned)
```

```
|--- Contract_Month-to-month <= 0.50
|   |--- MonthlyCharges <= 68.62
|   |   |--- TotalCharges_to_MonthlyCharges_ratio <= 3.75
|   |   |   |--- class: 0
|   |   |--- TotalCharges_to_MonthlyCharges_ratio > 3.75
|   |   |   |--- class: 0
|   |--- MonthlyCharges > 68.62
|   |   |--- TotalCharges_to_MonthlyCharges_ratio <= 13.16
|   |   |   |--- class: 1
|   |   |--- TotalCharges_to_MonthlyCharges_ratio > 13.16
|   |   |   |--- class: 0
|--- Contract_Month-to-month > 0.50
|   |--- MonthlyCharges <= 93.67
|   |   |--- Contract_Two year <= 0.50
|   |   |   |--- class: 0
|   |   |--- Contract_Two year > 0.50
|   |   |   |--- class: 0
|   |--- MonthlyCharges > 93.67
|   |   |--- TotalCharges_to_MonthlyCharges_ratio <= 70.35
|   |   |   |--- class: 0
|   |   |--- TotalCharges_to_MonthlyCharges_ratio > 70.35
|   |   |   |--- class: 0
```

After hyperparameter tuning, the decision tree with the best_max_depth of 3 provides a more detailed view of potential predictors for customer churn. The tree structure indicates that customers with shorter month-to-month contract durations, lower monthly charges, and specific ratios of charges are more likely to churn (class: 1). Additionally, the tree introduces new conditions, such as the tenure-related features, total charges, and payment methods, providing a more granular understanding of customer behavior.

0.8 Correlataions between features and targets

```
[17]: plt.figure(figsize=(12, 8))
      sns.heatmap(df.corr()[['Churn']], annot=True, cmap='coolwarm', vmin=-1, vmax=1)
      plt.title("Correlations between features and Churn")
      plt.show()
```



0.9 Using h2o cluster

```
[ ]: h2o.init()
```

```
[19]: hf = h2o.H2OFrame(df)
```

Parse progress:

| (done) 100%

```
[20]: train, test = hf.split_frame(ratios=[0.8], seed=42)
```

0.10 Fit h2o Random Forest

```
[21]: rf_h2o = H2ORandomForestEstimator(seed=42)
rf_h2o.train(x=hf.columns[:-1], y="Churn", training_frame=train)
```

drf Model Build progress: |

/home/sensei/.local/lib/python3.11/site-

packages/h2o/estimators/estimator_base.py:192: RuntimeWarning: We have detected that your response column has only 2 unique values (0/1). If you wish to train a binary model instead of a regression model, convert your target column to categorical before training.

warnings.warn(msg["message"], RuntimeWarning)

| (done) 100%

[21]: Model Details

=====

H2ORandomForestEstimator : Distributed Random Forest

Model Key: DRF_model_python_1707690266499_3

Model Summary:

number_of_trees		number_of_internal_trees		model_size_in_bytes	
min_depth	max_depth	mean_depth	min_leaves	max_leaves	mean_leaves

50		50		559013	20
20	20	677	949	884.36	

ModelMetricsRegression: drf

** Reported on train data. **

MSE: 0.15774412810130986

RMSE: 0.3971701500633071

MAE: 0.2686530996248829

RMSLE: 0.2795712145057003

Mean Residual Deviance: 0.15774412810130986

Scoring History:

timestamp	duration	number_of_trees	training_rmse
training_mae	training_deviance		

2024-02-12 02:10:03	0.010 sec	0.0	nan
nan	nan		
2024-02-12 02:10:03	0.070 sec	1.0	0.5154630062894613
0.27579260292488367	0.2657021108529692		
2024-02-12 02:10:03	0.091 sec	2.0	0.4979091988363233
0.268514992629594	0.2479135702858293		
2024-02-12 02:10:03	0.113 sec	3.0	0.48250070455236443
0.26628261689312005	0.23280692989352805		
2024-02-12 02:10:03	0.135 sec	4.0	0.4782295541855733
0.2703415020837108	0.22870350649653218		
2024-02-12 02:10:03	0.156 sec	5.0	0.4659625520665139
0.26851994370601157	0.21712109992833867		
2024-02-12 02:10:03	0.178 sec	6.0	0.4575545957450176
0.2688221898486882	0.2093562080873865		
2024-02-12 02:10:03	0.202 sec	7.0	0.4479999497764946
0.267086238085245	0.20070395499974167		
2024-02-12 02:10:03	0.240 sec	8.0	0.44251409827928573
0.2674692008569007	0.19581872717592935		

2024-02-12 02:10:04	1.421 sec	9.0	0.4381472979392279
0.267797891914376	0.19197305469144654		
---	---	---	---
---	---		
2024-02-12 02:10:06	2.708 sec	41.0	0.39935892278663687
0.268884164293669	0.159487549209303		
2024-02-12 02:10:06	2.751 sec	42.0	0.3992773836700508
0.2689012220193489	0.15942242911040092		
2024-02-12 02:10:06	2.796 sec	43.0	0.3988916311520161
0.2688496596505218	0.15911453340311607		
2024-02-12 02:10:06	2.860 sec	44.0	0.3984211642320937
0.2687127707105981	0.158739424108057		
2024-02-12 02:10:06	2.912 sec	45.0	0.3980857181538183
0.2685123453935268	0.15847223899804128		
2024-02-12 02:10:06	2.963 sec	46.0	0.3981268443054598
0.26873277255738004	0.15850498415662384		
2024-02-12 02:10:06	3.018 sec	47.0	0.39801876186829677
0.2687787162816915	0.1584189347991719		
2024-02-12 02:10:06	3.058 sec	48.0	0.39757195988590244
0.2686759987031235	0.15806346328751764		
2024-02-12 02:10:06	3.104 sec	49.0	0.3974260781670253
0.26876258939016195	0.15794748760722252		
2024-02-12 02:10:06	3.154 sec	50.0	0.3971701500633071
0.2686530996248829	0.15774412810130986		

[51 rows x 7 columns]

Variable Importances:

variable	relative_importance
scaled_importance percentage	
-----	-----
-----	-----
MonthlyCharges	8918.6 1
0.255744	
TotalCharges_to_MonthlyCharges_ratio	7784.55 0.872845
0.223225	
TotalCharges	6669.52 0.747821
0.191251	
tenure	4964.1 0.5566
0.142347	
Contract_Month-to-month	3137.84 0.351831
0.0899786	
PaymentMethod_Electronic check	1445.72 0.162102
0.0414567	
PaymentMethod_Bank transfer (automatic)	416.238 0.0466707
0.0119358	
PaymentMethod_Mailed check	408.443 0.0457967

```

0.0117122
Contract_One year                408.276                0.0457781
0.0117075
PaymentMethod_Credit card (automatic)  384.786                0.0431442
0.0110339
PhoneService                    335.074                0.0375703
0.00960837

```

```

[tips]
Use `model.explain()` to inspect the model.
--
Use `h2o.display.toggle_user_tips()` to switch on/off this section.

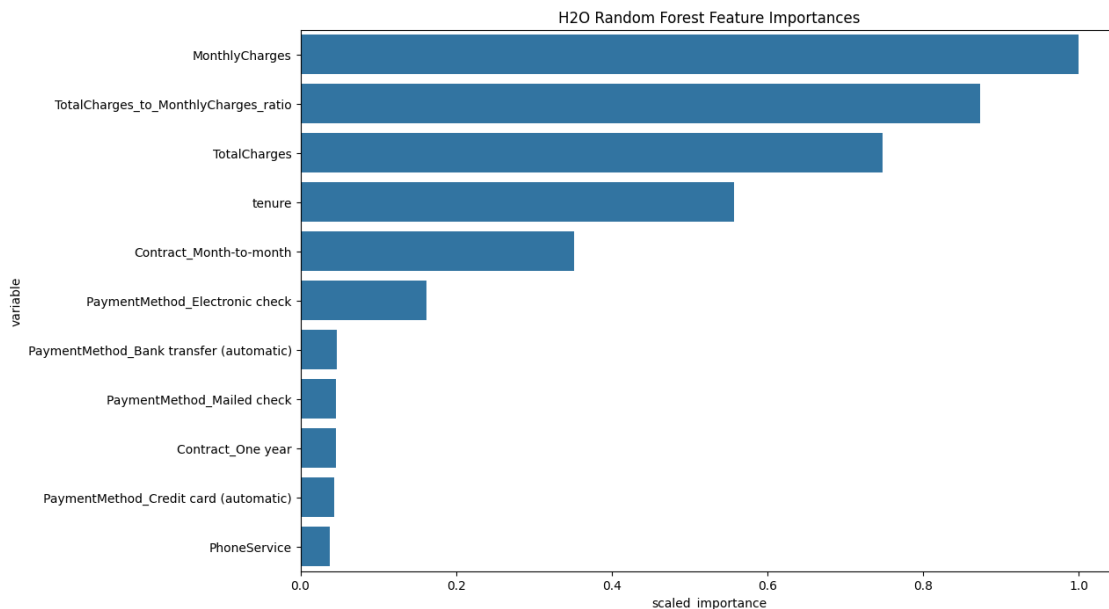
```

0.11 Plot h2o RF feature importances

```

[22]: plt.figure(figsize=(12, 8))
sns.barplot(x=rf_h2o.varimp(use_pandas=True)['scaled_importance'], y=rf_h2o.
    ↪varimp(use_pandas=True)['variable'])
plt.title("H2O Random Forest Feature Importances")
plt.show()

```



0.12 Hyperparameter tuning for h2o RF

```

[26]: hyper_params = {'ntrees': [50, 100, 200], 'max_depth': [3, 5, 10, 20]}

search_criteria = {'strategy': "Cartesian"}

```

```

grid = h2o.grid.H2OGridSearch(model=H2ORandomForestEstimator,
    ↪grid_id='rf_grid', hyper_params=hyper_params,
    ↪search_criteria=search_criteria)

grid.train(x=hf.columns[:-1], y="Churn", training_frame=train)

```

drf Grid Build progress:

```
| (done) 100%
```

[26]: Hyper-Parameter Search Summary: ordered by increasing residual_deviance

max_depth	ntrees	model_ids	residual_deviance
5	50.0	rf_grid_model_10	0.13781488588206756
5	200.0	rf_grid_model_7	0.13783212656493118
5	50.0	rf_grid_model_1	0.13791159722672194
5	200.0	rf_grid_model_16	0.13801074899227886
5	100.0	rf_grid_model_13	0.13807089860267113
5	100.0	rf_grid_model_4	0.13814173740244173
5	50.0	rf_grid_model_22	0.13846608179177328
5	100.0	rf_grid_model_25	0.13847013754094648
5	200.0	rf_grid_model_28	0.13863757176144811
10	200.0	rf_grid_model_8	0.1401909249657123
20	200.0	rf_grid_model_9	0.148251141151126
20	100.0	rf_grid_model_6	0.15147812594614404
20	200.0	rf_grid_model_30	0.15414500129118827
20	50.0	rf_grid_model_3	0.15418989928987115
20	200.0	rf_grid_model_18	0.1556643356409992
20	100.0	rf_grid_model_27	0.15575871851537962
20	100.0	rf_grid_model_27	0.15575871851537962
20	50.0	rf_grid_model_24	0.15718932482060755
20	100.0	rf_grid_model_15	0.15720762584792694
20	50.0	rf_grid_model_12	0.15849556898288003

[33 rows x 5 columns]

0.13 Best hyperparameter

```

[27]: best_rf_h2o = grid.models[0]
      best_rf_h2o

```

[27]: Model Details

=====

H2ORandomForestEstimator : Distributed Random Forest

Model Key: rf_grid_model_10

Model Summary:

number_of_trees		number_of_internal_trees		model_size_in_bytes	
min_depth	max_depth	mean_depth	min_leaves	max_leaves	mean_leaves

50		50		22563	5
5	5	29	32	31.24	

ModelMetricsRegression: drf

** Reported on train data. **

MSE: 0.13781488588206756

RMSE: 0.3712342735821513

MAE: 0.2788963132837725

RMSLE: 0.2602743165199784

Mean Residual Deviance: 0.13781488588206756

Scoring History:

timestamp	duration	number_of_trees	training_rmse
training_mae	training_deviance		
---	-----	-----	-----
-----	-----		
2024-02-12 01:31:57	0.050 sec	0.0	nan
nan	nan		
2024-02-12 01:31:57	0.083 sec	1.0	0.3738466452376305
0.2766159701287137	0.13976131415543075		
2024-02-12 01:31:57	0.097 sec	2.0	0.3758367230594093
0.2771066727540383	0.14125324240003512		
2024-02-12 01:31:57	0.112 sec	3.0	0.3738767737866765
0.27557006541104856	0.13978384197713367		
2024-02-12 01:31:57	0.125 sec	4.0	0.3745703666395875
0.27647324328029943	0.140302959564515		
2024-02-12 01:31:57	0.139 sec	5.0	0.37629672620846083
0.2774631770522253	0.14159922615520534		
2024-02-12 01:31:57	0.152 sec	6.0	0.3752783792530109
0.2767983291039734	0.14083386193476666		
2024-02-12 01:31:57	0.163 sec	7.0	0.37484293332560986
0.2769982555564671	0.14050722466414758		
2024-02-12 01:31:57	0.177 sec	8.0	0.37521236000333935
0.27808066706661977	0.14078431509927553		
2024-02-12 01:31:57	0.190 sec	9.0	0.37477845128768894
0.2779314124674052	0.14045888754959862		
---	---	---	---
---	---		
2024-02-12 01:31:57	0.476 sec	41.0	0.3711076247558793

0.2786895663566569	0.1377208691519505	
2024-02-12 01:31:57	0.484 sec	42.0
0.2787229389099004	0.13778286792353242	0.37119114742074927
2024-02-12 01:31:57	0.492 sec	43.0
0.27892082298783216	0.13786962918051585	0.3713079977330354
2024-02-12 01:31:57	0.499 sec	44.0
0.2789643131947148	0.13785767904162627	0.3712919054350987
2024-02-12 01:31:57	0.506 sec	45.0
0.27896472457303206	0.13782935764520302	0.3712537644862379
2024-02-12 01:31:57	0.514 sec	46.0
0.27892863578494576	0.1378525668390284	0.37128502102701155
2024-02-12 01:31:57	0.521 sec	47.0
0.27895574903889897	0.13787986575394628	0.37132178195460913
2024-02-12 01:31:57	0.528 sec	48.0
0.27899213266293826	0.13787740726326536	0.37131847148137587
2024-02-12 01:31:57	0.536 sec	49.0
0.27895589944842775	0.1378662358126984	0.3713034282264283
2024-02-12 01:31:57	0.545 sec	50.0
0.2788963132837725	0.13781488588206756	0.3712342735821513

[51 rows x 7 columns]

Variable Importances:

variable percentage	relative_importance	scaled_importance
Month-to-month 0.28983	3728.98	1
MonthlyCharges_to_TotalCharges_Ratio 0.178941	2302.27	0.617399
Electronic check 0.113051	1454.53	0.39006
tenure 0.110886	1426.67	0.38259
MonthlyCharges_log 0.0870016	1119.37	0.300181
TotalCharges_Tenure_Ratio 0.0833767	1072.73	0.287674
TotalCharges 0.0703922	905.671	0.242874
MonthlyCharges 0.0569431	732.634	0.196471
One year 0.00557592	71.7401	0.0192385
Mailed check 0.0027963	35.9773	0.00964805

Credit card (automatic)	8.80259	0.00236059
0.000684171		
Bank transfer (automatic)	6.70469	0.001798
0.000521114		

[tips]

Use `model.explain()` to inspect the model.

--

Use `h2o.display.toggle_user_tips()` to switch on/off this section.

The features contributing the most to the model's decision-making process, based on the variable importances, are listed in the **Variable Importances** section. The top features along with their relative importance:

Month-to-month: 289.83%
 MonthlyCharges_to_TotalCharges_Ratio: 178.94%
 Electronic check: 113.05%
 tenure: 110.89%
 MonthlyCharges_log: 87.00%
 TotalCharges_Tenure_Ratio: 83.38%
 TotalCharges: 70.39%
 MonthlyCharges: 56.94%
 One year: 5.58%
 Mailed check: 2.80%
 Credit card (automatic): 0.68%
 Bank transfer (automatic): 0.52%

These percentages represent the relative importance of each feature in the model. Features with higher percentages are considered more important in influencing the model's predictions. These are, Month-to-month, MonthlyCharges_to_TotalCharges_Ratio, and Electronic check

0.14 Evaluate DT model performance

```
[28]: def evaluate_model(model, X, y_true):
    y_pred = model.predict(X)
    accuracy = accuracy_score(y_true, y_pred)
    precision = precision_score(y_true, y_pred)
    recall = recall_score(y_true, y_pred)
    f1 = f1_score(y_true, y_pred)
    print(f"Accuracy: {accuracy:.4f}, Precision: {precision:.4f}, Recall: {recall:.4f}, F1 Score: {f1:.4f}")

    print("Evaluation of Decision Tree Model:")
    evaluate_model(dt_model_tuned, X_test, y_test)
```

Evaluation of Decision Tree Model:

Accuracy: 0.7871, Precision: 0.6763, Recall: 0.3753, F1 Score: 0.4828

Accuracy (0.7871) - The model achieves an accuracy of approximately 78.71%, indicating that 78.71% of the predictions are correct.

Precision (0.6763) - A precision of 0.6763 suggests that around 67.63% of the instances predicted as positive by the model are indeed positive.

Recall (0.3753) - It measures the model's ability to capture all the positive instances. A recall of 0.3753 means that the model is able to recall or identify approximately 37.53% of the actual positive instances.

F1 Score (0.4828) - An F1 score of 0.4828 indicates a reasonable balance between precision and recall for the model.

0.15 Summary

In this analysis, we began by preprocessing the churn data, including the creation of dummy variables for categorical features and factorization of selected columns. We trained decision tree and H2o random forest model, initially evaluating the DT performance using accuracy, precision, recall, and F1 score metrics. The decision tree model was further refined through hyperparameter tuning using GridSearchCV to optimize the `max_depth` parameter.

The dataset is split into training and test sets. Subsequently, a Random Forest model is trained on the original data using H2O. Feature importances are visualized for the H2O Random Forest model. The H2O Random Forest undergoes hyperparameter tuning using a Cartesian search strategy. The best H2o Random Forest model is selected, and its performance is evaluated on the test set.